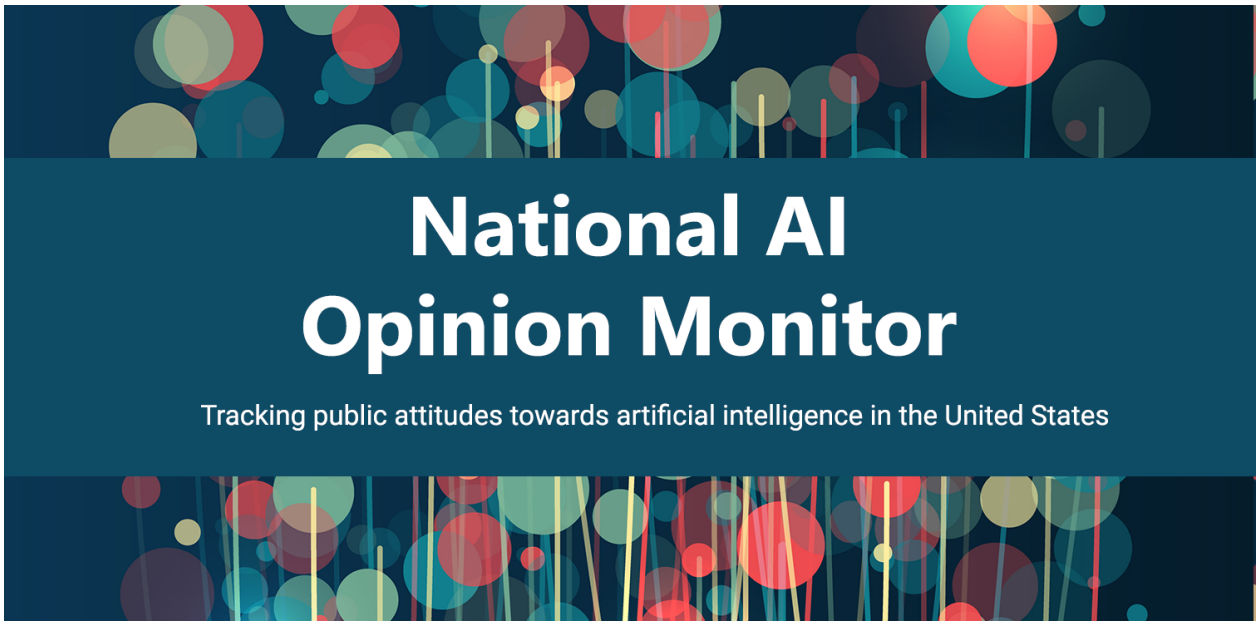

www.naiom.net

NATIONAL AI OPINION MONITOR: AI TRUST AND KNOWLEDGE IN AMERICA

USA, February 2025

Katherine Ognyanova, Rutgers University
Vivek Singh, Rutgers University

NAIOM
National AI Opinion Monitor

RUTGERS UNIVERSITY

COVER MEMO

Report of February 1, 2025, v.1

National AI Opinion Monitor

Rutgers University

School for Communication & Information

NAIOM
National AI Opinion Monitor

Authors:
Katherine Ognyanova, Rutgers University
Vivek Singh, Rutgers University

Note on methods:

Between 10/25/2024, and 11/08/2024, we collected a total of 4,767 responses from participants ages 18 and older living in the US. The surveys were conducted by PureSpectrum via an online, nonprobability sample, with representative interlocked quotas for gender, age, race, ethnicity, and region. In addition to balancing on these dimensions, we reweighted our data to match the U.S. population with respect to gender, age, race, ethnicity, education, region, internet use, and living in urban, suburban, or rural areas. For more information, visit www.naiom.net.

Contact information:

For additional information and press requests, contact:

- Katherine Ognyanova at katya.ognyanova@rutgers.edu
- Vivek Singh at v.singh@rutgers.edu

AI trust and knowledge in America

Artificial intelligence (AI) has quickly become part of our daily routine. We are increasingly exposed to these technologies, both at work and in our personal lives. The National AI Opinion Monitor (NAIOM) tracks public perceptions about AI, offering a detailed look at attitudes across demographic groups. The insights we offer come at a crucial time: as AI continues to expand into media, healthcare, law enforcement, education, and beyond, understanding public sentiment is essential for policymakers, companies, journalists, educators, and the general public.

In a previous report, we examined the awareness and use of AI tools among the US public, as well as the key hopes and concerns of Americans with regard to artificial intelligence.

Here, we examine public trust in AI, in the companies that use it, and in the news content produced by it. We also evaluate people's objective and perceived knowledge about artificial intelligence. The next report in this series will focus on AI, the future of work, and attitudes towards regulation.

KEY TAKEAWAYS

- Our results demonstrate an emerging AI divide in the US. People with higher income and education are more likely not only to use AI, but also to trust it, as well as to have higher objective and perceived knowledge about it.
- Close to half (47%) of Americans report "a fair amount" or "a great deal" of confidence in AI to act in the public interest. Confidence is higher among men (52%), non-White respondents (55%), those in age group 25-44 (55%), graduate degree holders (60%), high-income earners ($100K+, 63%), Democrats (56%), and urban area residents (53%).
- Half of Americans report they trust companies "some" or "a lot" to use AI responsibly. Trust in companies follows similar demographic patterns to confidence in AI. It is most prevalent among those with higher socioeconomic status: graduate degree holders (65%) and high-income earners ($100K+, 65%).

- Americans across all demographic groups trust information produced by mainstream media journalists more than AI-generated information. Overall, 62% trust journalistic content “some” or “a lot”, compared to 48% who trust AI content. Even so, public trust in AI information is quite high, especially among Asian Americans (62%), graduate degree holders (64%), and those with high income ($100K+, 64%).

- About 43% of respondents reported being “somewhat” or “very” confident that they could recognize whether a news story was written by a human or AI. Confidence levels were highest among people aged 18-24 (60%), graduate degree holders (58%), and high-income Americans ($100K+, 62%).

- Respondents reported only slightly higher average generalized trust in humans compared to their trust in machines (5.4 vs. 5.1 on a 10-point scale).

- Americans who had the highest and the lowest objective knowledge about artificial intelligence were less likely to trust AI (38%), while those with medium knowledge levels were more trusting (57%).

- We measured objective knowledge of AI using an 8-item true/false test. Only 23% of respondents were able to correctly identify more than half of the items, while less than 1% gave correct answers for all 8 items. Higher income and education were strongly associated with greater knowledge.

- About a quarter (26%) of Americans reported hearing “a lot” about AI, while 63% had heard “a little” and 11% “nothing at all.” Self-reported familiarity was highest among men (33%), urban residents (34%), and those with higher education and income.

American trust in artificial intelligence

Trust in artificial intelligence is crucial for its successful integration into our society. While AI is quickly becoming an increasingly important part of our work, education, and public life, its adoption and use are still premised on public trust. Trust in AI is a multifaceted concept that depends on context— people's reported attitudes may differ when it comes to trusting the technology itself, trusting the companies behind it, or trusting the information it produces. In this report, we ask about all of these attitudes and report on the demographic patterns of trust across all of them. In the second part of the report, we review the objective and perceived AI knowledge among Americans.

Asked about their trust in AI to act in the public interest, 47% of Americans reported having "a fair amount" or "a great deal" of confidence, a percentage higher than that for social media (39%) or Congress (42%). Men were more likely to trust AI (Figure 1), with 52% having confidence in it, compared to 43% of women. Education also emerged as a significant factor: graduate degree holders had most confidence in AI (60% reporting "a fair amount" or "a great deal"), while those without a college degree were less likely to be confident (45%). There were predictable age gaps, with younger respondents more likely to report high confidence (47% of those aged 18-24 and 55% of those aged 25-44), compared to only 39% of those 65 and older. Racial and ethnic differences were also notable: Asian American (58%), African American (57%), and Hispanic (55%) respondents reported higher confidence compared to Whites (43%). Trust levels differed across political lines, with 56% of Democrats expressing "a fair amount" or "a great deal" of confidence, compared to 46% of Republicans and 40% of Independents. Geographically, trust was higher among those living in urban areas (53%) compared to suburban (47%) and rural (38%) areas. In models where all demographic variables are used to simultaneously predict confidence in AI, the most important predictors were being male, being African American, lower age, higher income, and being a Democrat.

The public's trust in companies to use AI responsibly reflected similar patterns. Half of the respondents reported trusting companies "some" or "a lot" (Figure 2). Gender differences remained consistent, with 54% of men selecting the top two categories for trust, compared to 46% of women. Graduate degree holders reported higher trust in companies, with 65% selecting "some" or "a lot", compared to only 46% of those with no college degree. Income was also linked to trust: respondents earning over $100K reported 65% in the highest two categories, compared to only 45% of those earning under $25K. Geographically, urban residents were more likely to express trust (53%) than rural respondents (42%). In models examining all demographic factors together, being male, a Democrat, having higher education and income predicted more trust in companies to use AI responsibly.

Trust in artificial intelligence to act in the public interest

How much confidence, if any, do you have in each of the following, to act in the best interest of the public? - Artificial Intelligence (AI)

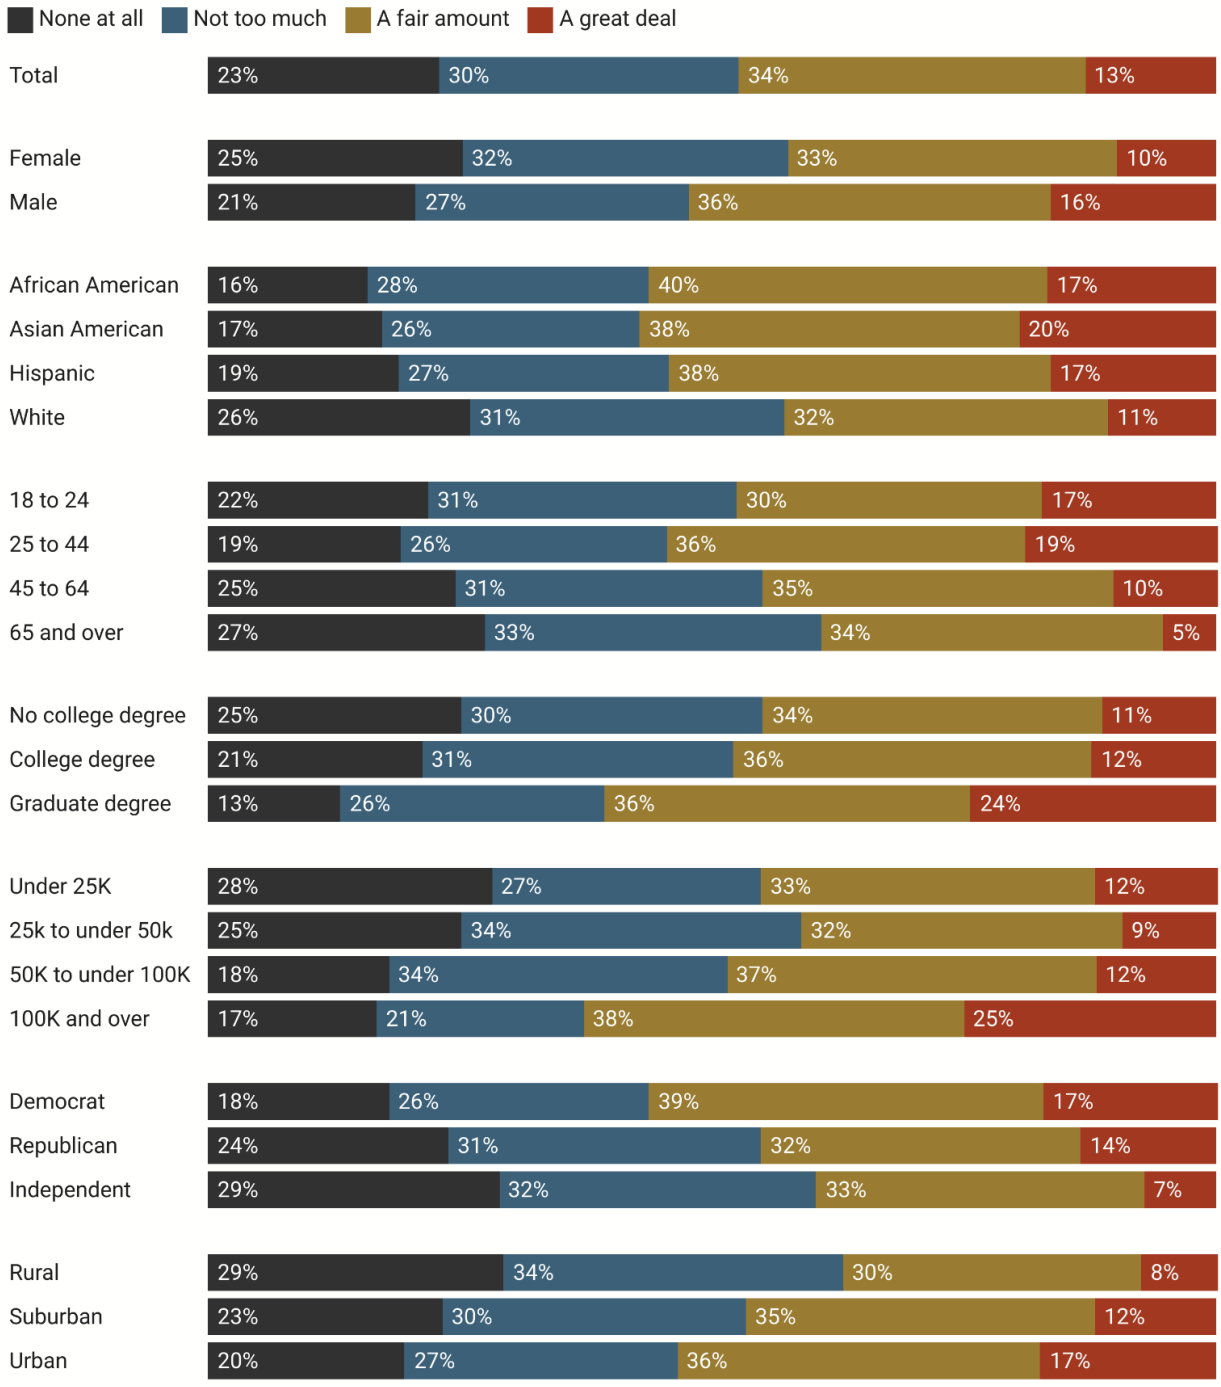

National AI Opinion Monitor (NAIOM). National sample, December 2024, N = 4767. More information at www.naiom.net.

Figure 1.

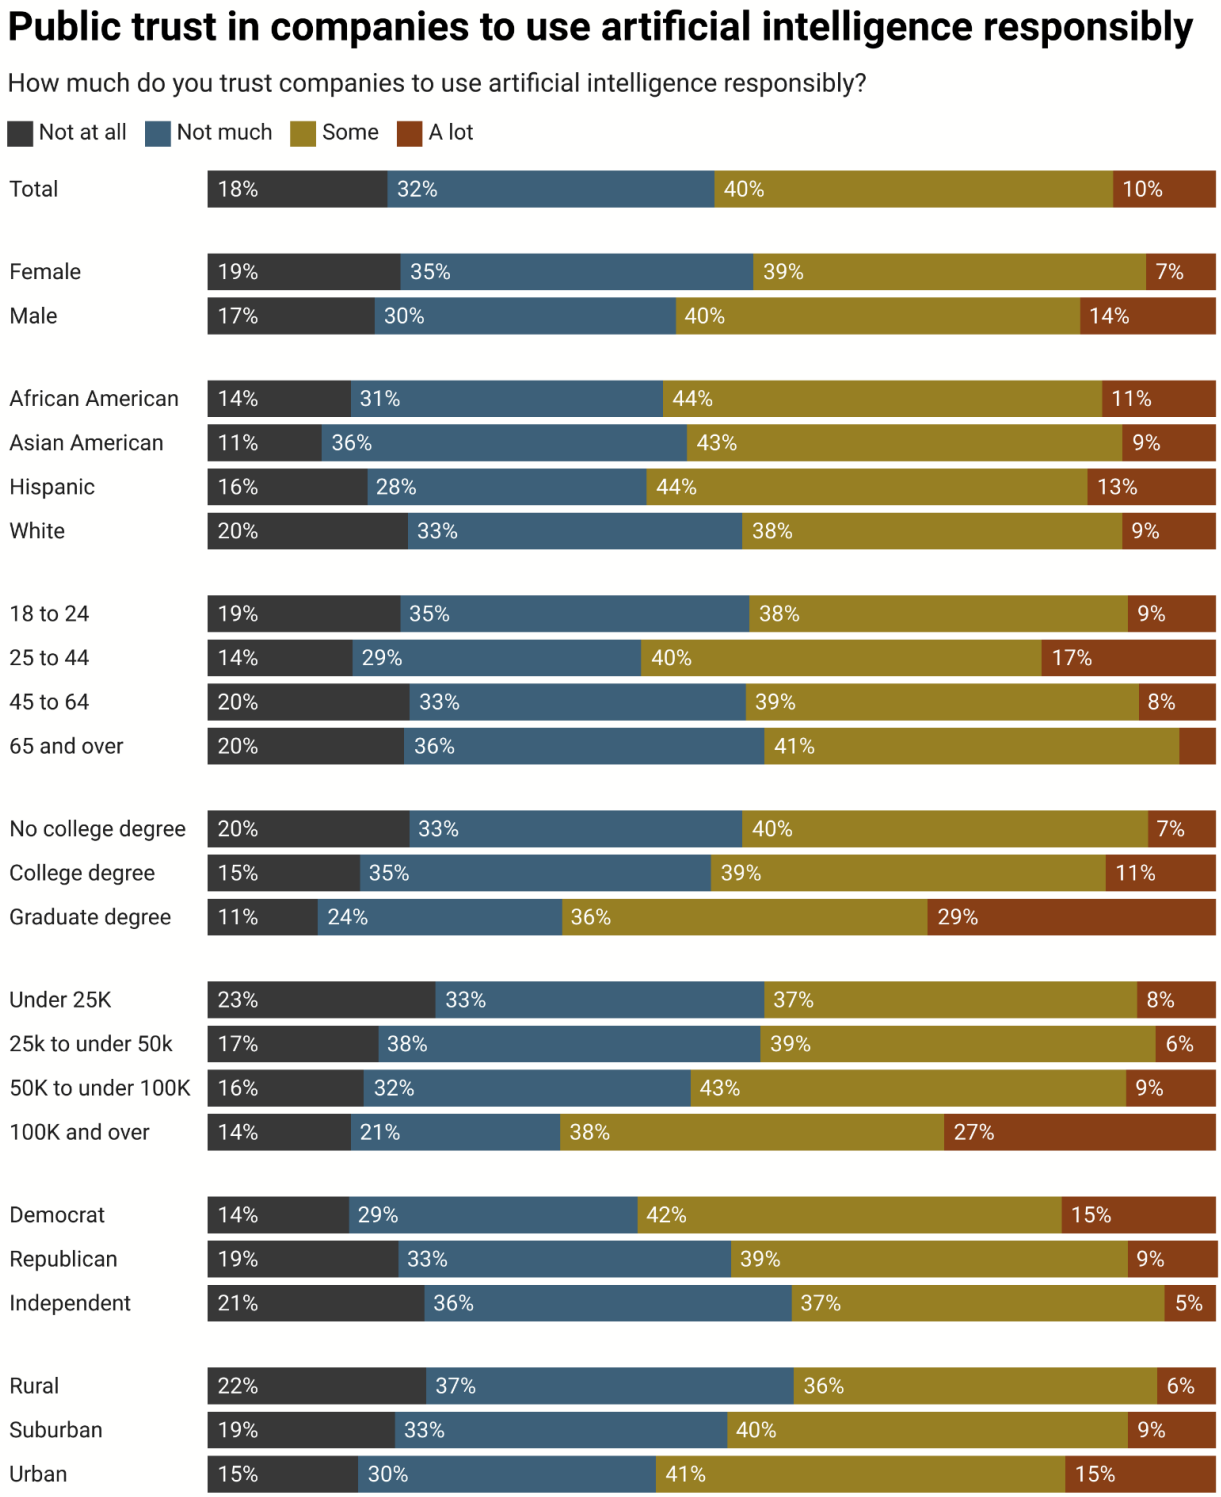
Public trust in companies to use artificial intelligence responsibly
How much do you trust companies to use artificial intelligence responsibly?
Not at all
Not much
Some
A lot
Total 18% 32% 40% 10%
Female 19% 35% 39% 7%
Male 17% 30% 40% 14%
African American 14% 31% 44% 11%
Asian American 11% 36% 43% 9%
Hispanic 16% 28% 44% 13%
White 20% 33% 38% 9%
18 to 24 19% 35% 38% 9%
25 to 44 14% 29% 40% 17%
45 to 64 20% 33% 39% 8%
65 and over 20% 36% 41%
No college degree 20% 33% 40% 7%
College degree 15% 35% 39% 11%
Graduate degree 11% 24% 36% 29%
Under 25K 23% 33% 37% 8%
25k to under 50k 17% 38% 39% 6%
50K to under 100K 16% 32% 43% 9%
100K and over 14% 21% 38% 27%
Democrat 14% 29% 42% 15%
Republican 19% 33% 39% 9%
Independent 21% 36% 37% 5%
Rural 22% 37% 36% 6%
Suburban 19% 33% 40% 9%
Urban 15% 30% 41% 15%
National AI Opinion Monitor (NAIOM). National sample, December 2024, N = 4767. More information at www.naiom.net.

Figure 2.

To provide some context, we also examined how much Americans tend to trust other humans, and how much they trust machines (Figure 3). We used a standard survey question measuring generalized social trust: "Generally speaking, would you say that most **people** can be trusted, or that you can't be too careful in dealing with people?" Respondents answered on a scale from 1 to 10, where 1 means "*you can't be too careful*" and 10 means "*most people can be trusted*". We also asked a second question, replacing "people" with "machines": "Generally speaking, would you say that most **machines** can be trusted, or that you can't be too careful in dealing with machines?".

Respondents reported slightly higher average trust in people (5.4) compared to machines (5.1). This pattern holds across most demographic groups, though the magnitude of the difference varies. While most differences are statistically significant, many are substantively small.

Compared to women, men reported higher levels of trust in both people and machines. On average, men scored 5.6 for trust in people, compared to a 5.1 average among women. For trust in machines, men scored 5.5, while the score was 4.8 for women. This is in line with findings from previous research, which has found lower trust scores among women, potentially due to higher risk factors or more risk aversion.

Most age groups reported relatively similar average trust in humans and machines (within 1 decimal point), with the exception of those aged 65+, who reported the lowest level of trust in machines (4.2) and highest level of trust in people (5.7).

Socioeconomic status was closely linked to trust, with higher income and education associated with more trust in both humans and machines. For both trust scores, there was over a point difference in the average scores of the people with lowest income and education compared to those with high earnings and graduate degrees. Interestingly, however, higher education was also linked to more divergent trust scores: graduate degree holders trusted people .5 more than machines, while the difference was .2 for those without a degree.

Differences also emerged with regard to politics. Democrats reported the highest trust in people (5.9) and machines (5.6), followed by Republicans (5.6 for people and 5.2 for machines). Independents had the least trust in both categories, with scores of 5.2 for people and 4.9 for machines.

Trust levels further varied by residential area. Urban residents reported the highest trust in people (5.7) and machines (5.6). Suburban respondents followed, with scores of 5.3 for people and 5.3 for machines. Rural respondents reported the lowest trust, with scores of 4.7 for people and 4.8 for machines.

Generalized trust in people and machines (1-10)

Generally speaking, would you say that most [people/machines] can be trusted, or that you can't be too careful in dealing with [people/machines] ?

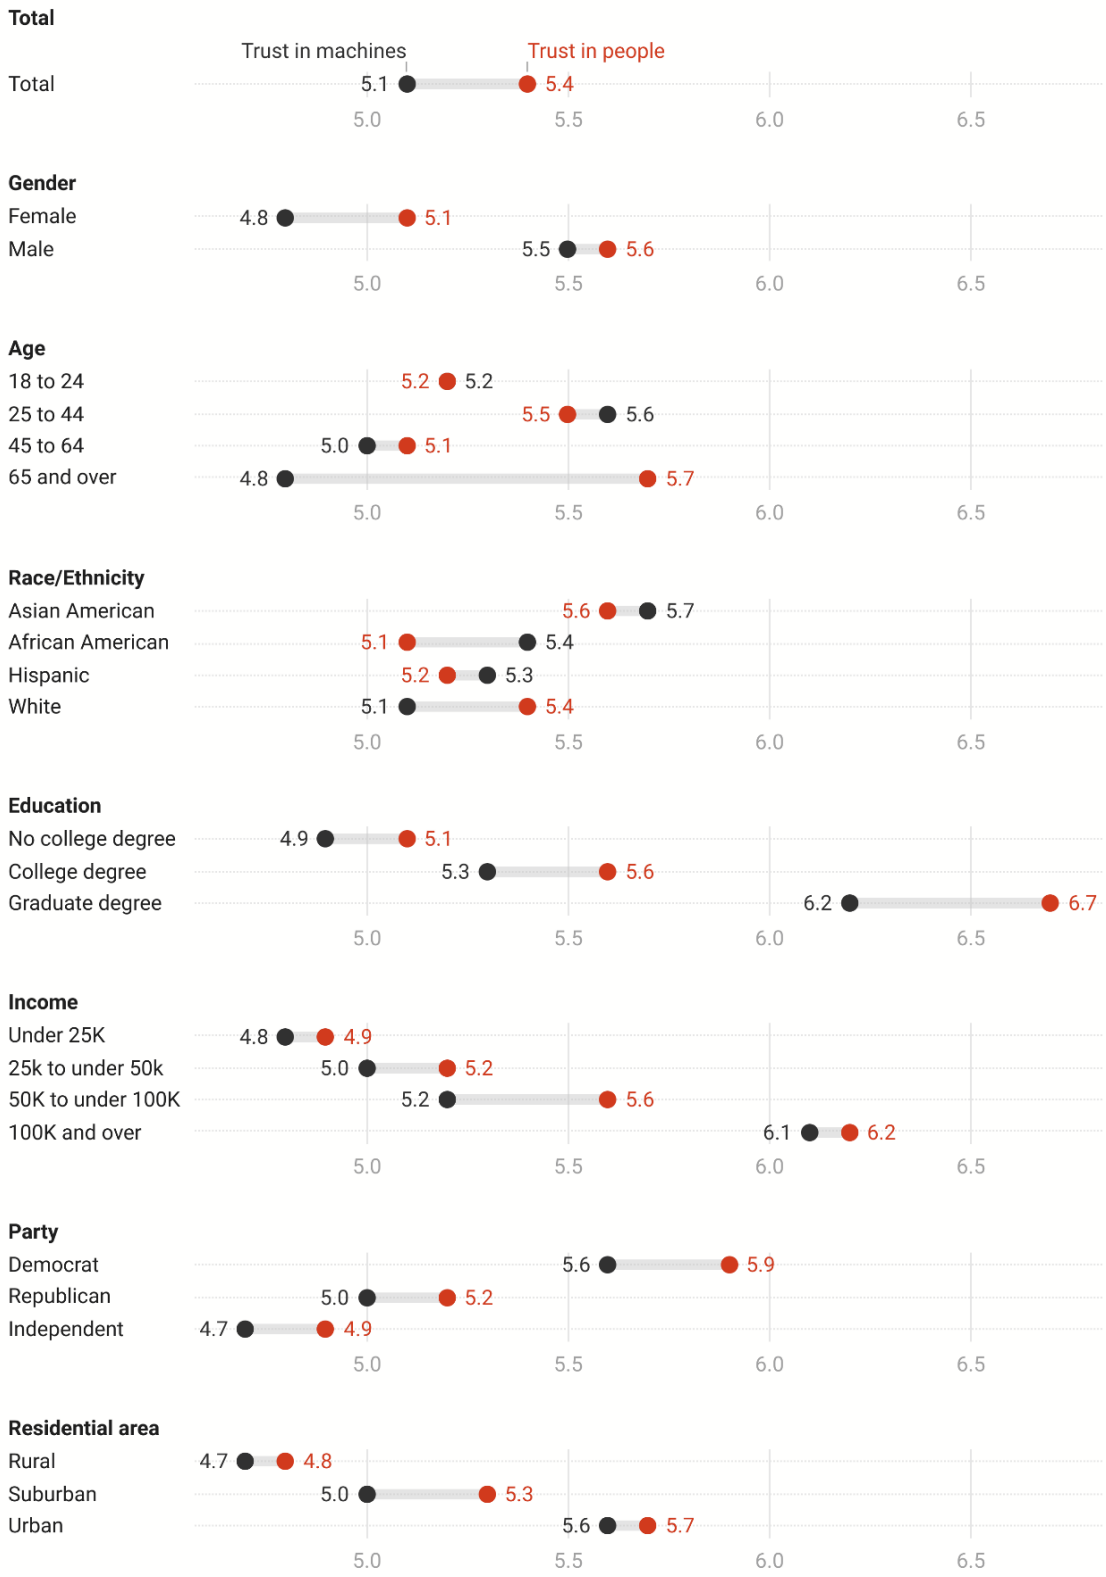

National AI Opinion Monitor (NAIOM). National sample, December 2024, N = 4767. More information at www.naiom.net.

Figure 3.

We further examined how much Americans trust information produced by AI, compared to their trust in information produced by journalists from mainstream media sources. Overall, respondents were more likely to trust news media compared to AI. In total, 62% of respondents reported trusting ("some" or "a lot") the information received from journalists, compared to 48% expressing trust in information from AI. Every single demographic group examined was more likely to trust news media materials compared to AI content (Figure 4).

Trust in information from journalists was higher among men (65%) compared to women (59%). Men also showed more trust in AI-generated information, with 54% reporting trust compared to 43% of women.

Racial and ethnic differences were also notable. African Americans (70%) and Asian Americans (67%) reported the highest levels of trust in information from journalists, with Hispanic respondents slightly lower at 63% and Whites at 59%. Trust in AI-generated information followed a similar trend, with Asian Americans (62%) leading, followed by African Americans (54%), Hispanics (53%), and Whites (45%).

Trust in AI was highest among respondents aged 25-44 (58%), and lowest among those aged 65 and older (40%). Information from journalists was most trusted among young respondents, reaching 67% for those aged 18-24.

Education and income were especially predictive of trust. Respondents with graduate degrees exhibited the highest trust in both journalists (76%) and AI (64%). Those with no college degree had lower trust levels, at 59% for journalists and 46% for AI. Trust in information from journalists and AI also increased with income. Respondents earning $100K or more expressed the highest levels of trust in journalists (69%) and AI (64%). In contrast, respondents earning under $25K exhibited the lowest trust levels, at 58% for journalists and 44% for AI.

Trust patterns also differed by political affiliation. Democrats reported the highest trust in journalists (78%) and AI (56%). In comparison, Republicans exhibited lower trust levels, with 48% trusting journalists and 46% trusting AI. Independents reported the lowest trust in AI-generated information (42%).

In terms of residential areas, urban respondents displayed the highest trust in both journalists (66%) and AI (54%), while rural residents had the lowest trust levels, at 53% for journalists and 39% for AI. Suburban respondents fell in between, with 62% trusting journalists and 48% trusting AI.

Overall, while trust is still higher for information from mainstream media, we also find very high levels of trust in AI content, especially among those with the highest education and income.

Trust in information produced by journalists vs. AI information

How much would you trust news or information produced by journalists from mainstream media?
How much would you trust news or information that was written by AI (artificial intelligence)?
[Percent respondents who trust the information "some" or "a lot"]

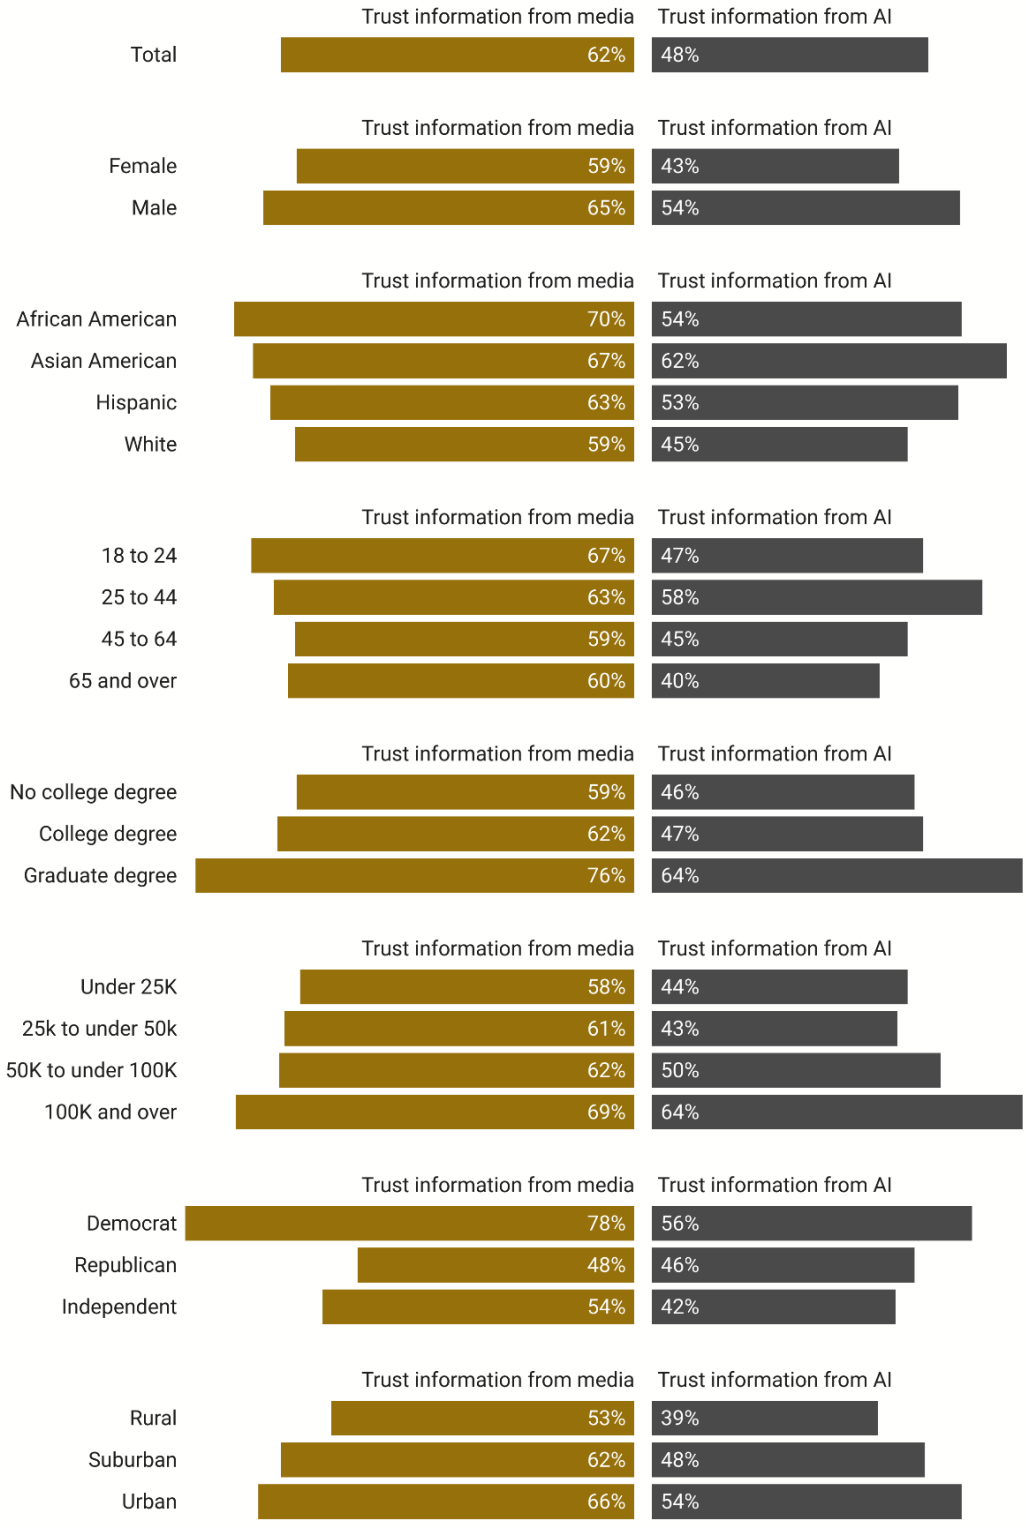

National AI Opinion Monitor (NAIOM). National sample, December 2024, N = 4767. More information at www.naiom.net.

Figure 4.

While people report different levels of trust in information from journalists and AI, they may not always be able to tell the difference between the two. Overconfidence in one's ability to differentiate might also be harmful. We asked respondents how confident they felt that they could recognize whether news stories were written by people or AI. Only a minority (13%) reported being very confident, while 30% said they were somewhat confident.

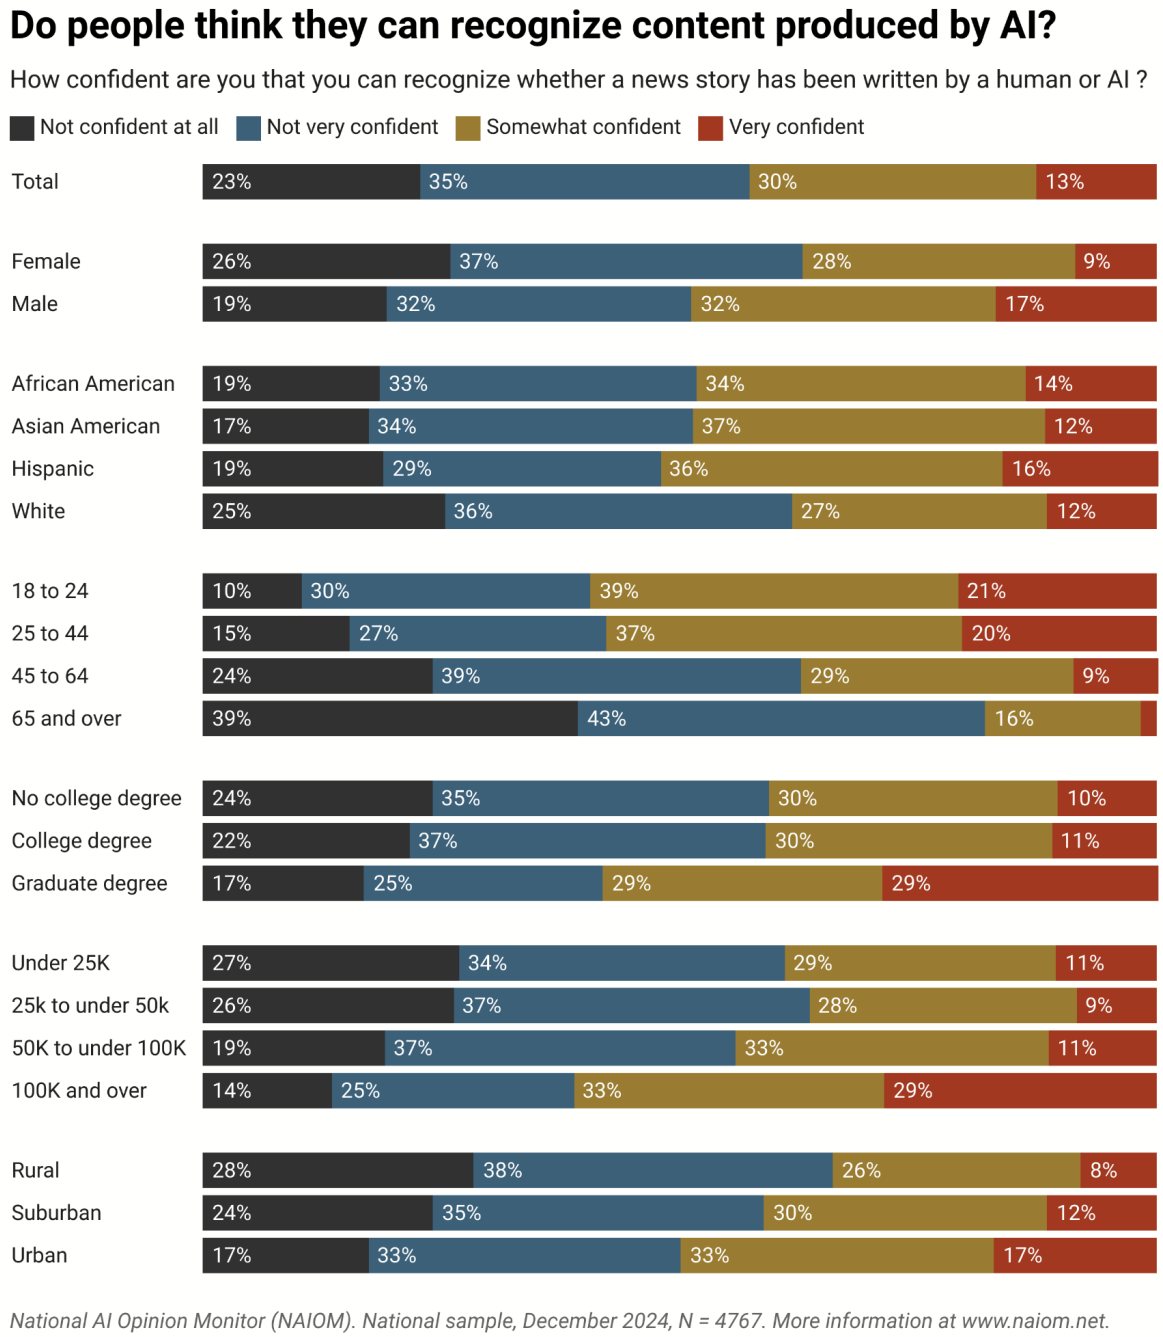

Figure 5.

Men were more likely to claim they could recognize AI news, with 49% reporting they were somewhat or very confident, compared to 37% of women. Confidence decreased with age. Over 60% of those aged 18-24 reported they were somewhat or very confident, compared to only 25% of respondents aged 65 and over. Hispanic Americans were most confident (52%), compared to 39% among Whites. Predictably, confidence was especially high among respondents with the highest levels of education (58%) and income (62%).

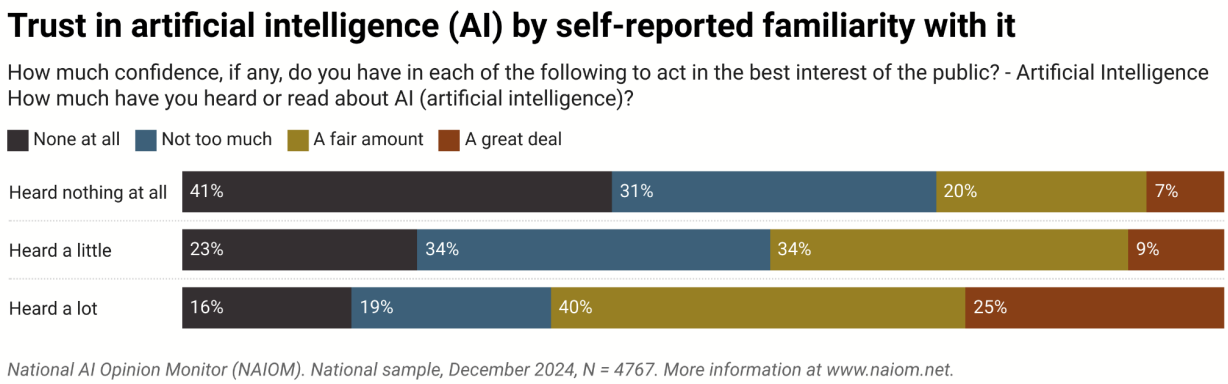

Figure 6.

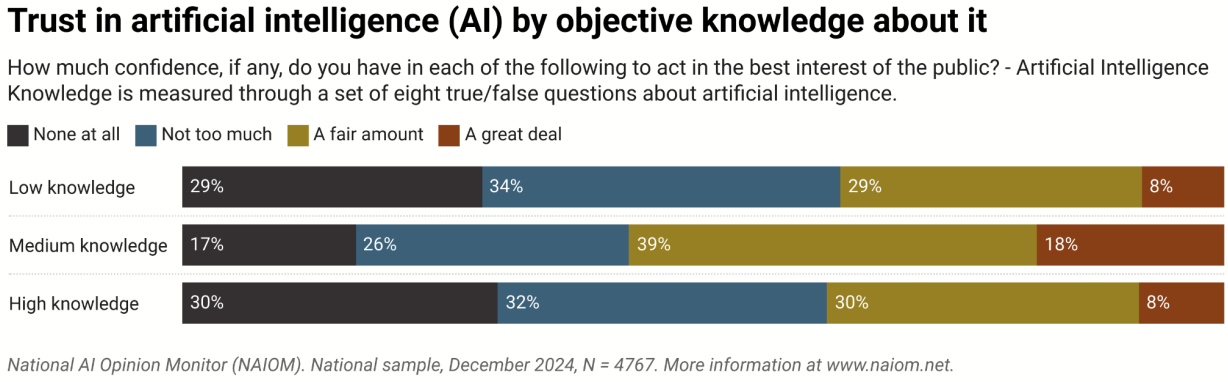

Figure 7.

In the next section of this report, we discuss in more detail the self-reported and objective familiarity Americans have with artificial intelligence. Here, we report how trust in artificial intelligence is linked to self-reported and objective knowledge about AI. The patterns we find differ in informative ways. For self-reported knowledge (Figure 6) , higher perceived familiarity with AI is linked to higher trust. Those who claim to have heard or read a lot about AI have higher trust levels than those who have not. When measuring knowledge objectively by asking respondents a series of true/false questions, we find that trust is highest among those with medium levels of knowledge (Figure 7). Respondents who objectively score lowest or highest on AI knowledge are less likely to trust AI. About 38% of people in those two categories trust AI “a fair amount” or “a great deal”, compared to 57% of those in the middle knowledge category.

Objective and self-reported knowledge about AI

To get a sense of how well people understand artificial intelligence, we presented them with eight statements (Figure 8) and asked them to classify each as “accurate”, “inaccurate”, or say they were “not sure”. Three of the statements were correct, the rest were not. We calculated respondent knowledge scores as the number of statements they correctly identified as being accurate or inaccurate. On average, respondents gave 3.3 correct answers, with 10% getting none of the answers right, and less than 1% getting all the answers correctly.

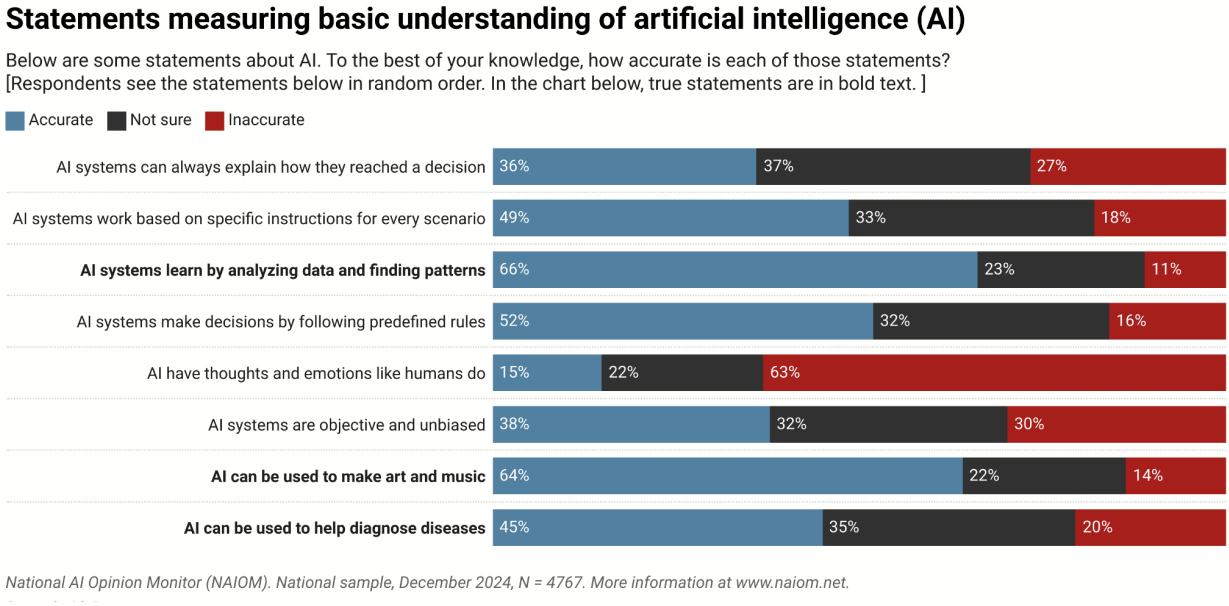

Figure 8.

Respondents were categorized into three groups: “low knowledge” (score 0-2), “medium knowledge” (score 3-4), and “high knowledge” (score 5-8). Overall, 27% of respondents fell into the low knowledge category, 51% into medium knowledge, and 23% into high knowledge.

Men were slightly more likely than women to fall into the high knowledge category, with 25% of men classified as high knowledge compared to 20% of women (Figure 9). Among racial and ethnic groups, Asian Americans had the highest percentage in the high knowledge category (23%), followed by Whites (26%). Black (18%) and Hispanic Americans (15%) were less likely to fall into this category. Younger respondents aged 18-24 had a relatively high representation in the medium knowledge group (52%) but lower representation in the high knowledge group (20%). Conversely, those aged 65 and over were more evenly distributed, with 49% in the medium category and 27% in the high category. Education had a substantial association with objective knowledge, as 29% of graduate degree holders were classified as high knowledge, compared to only 20% of

those without a college degree. Higher income was also linked to knowledge: 27% of respondents earning over $100K were in the high knowledge category, compared to 19% of those earning under $25K.).

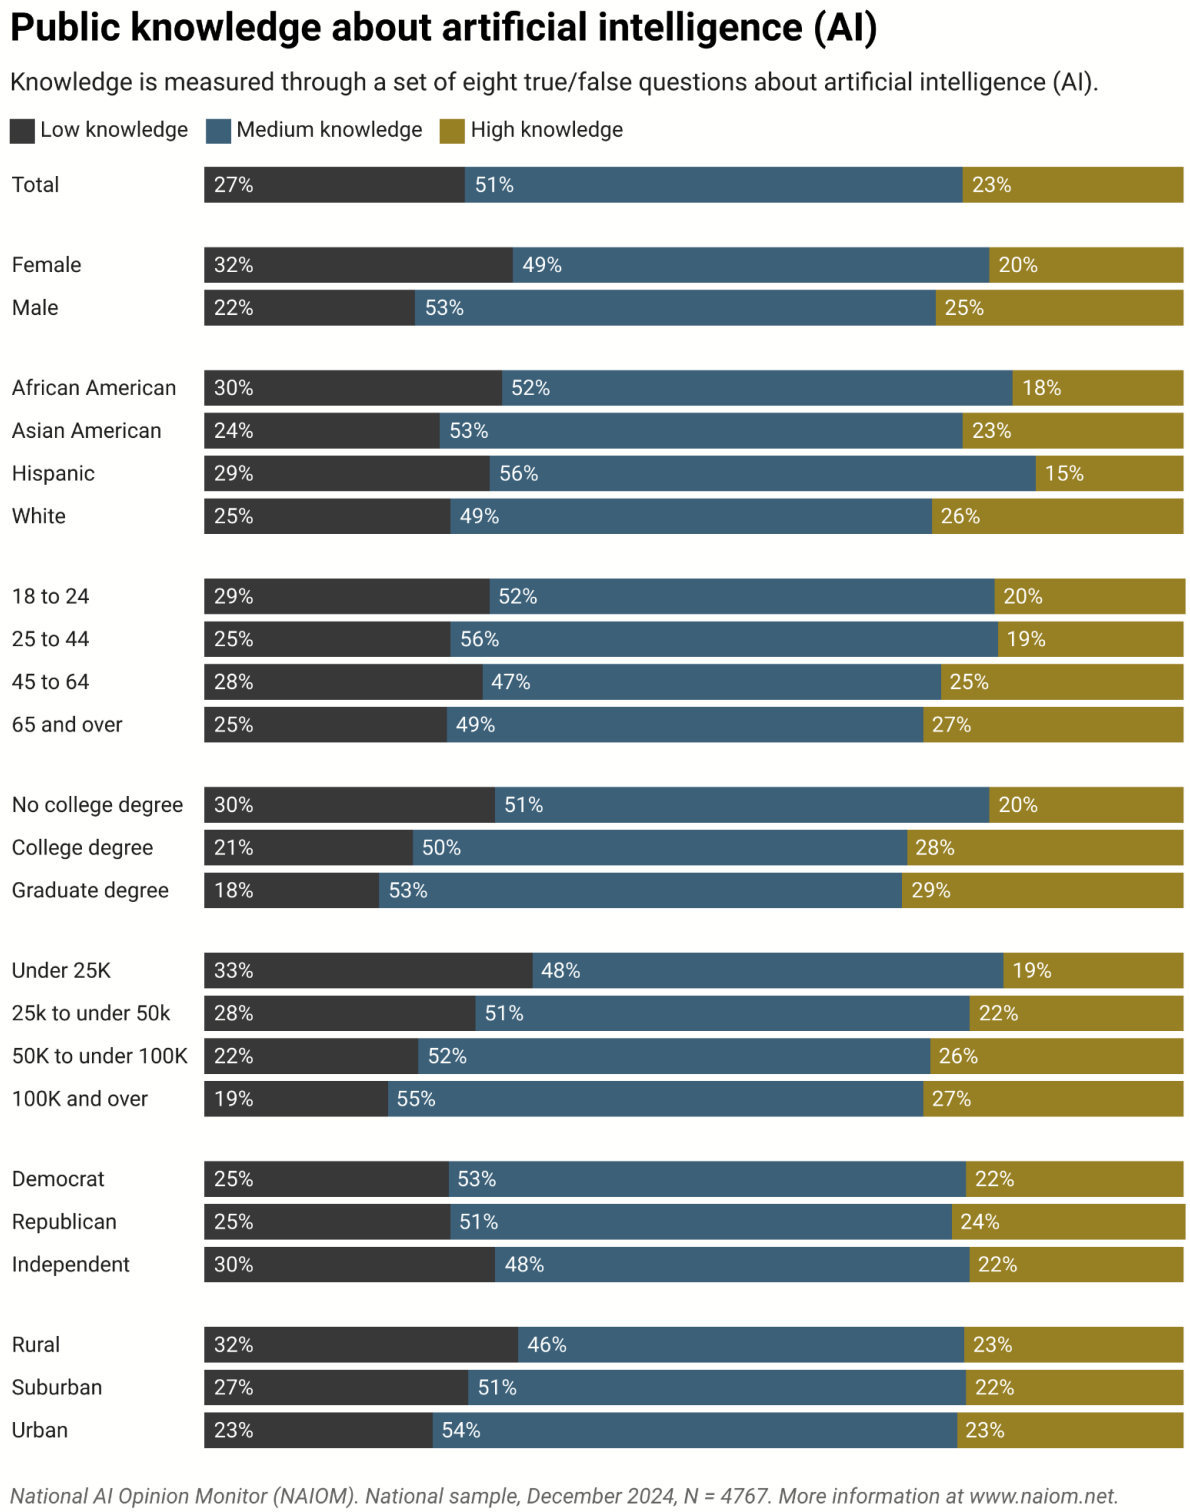

Figure 9.

Average knowledge about artificial intelligence (range 0-8)

Knowledge is measured through a set of eight true/false questions about artificial intelligence (AI).
[Red dots represent mean values, gray ranges are 95% confidence intervals]

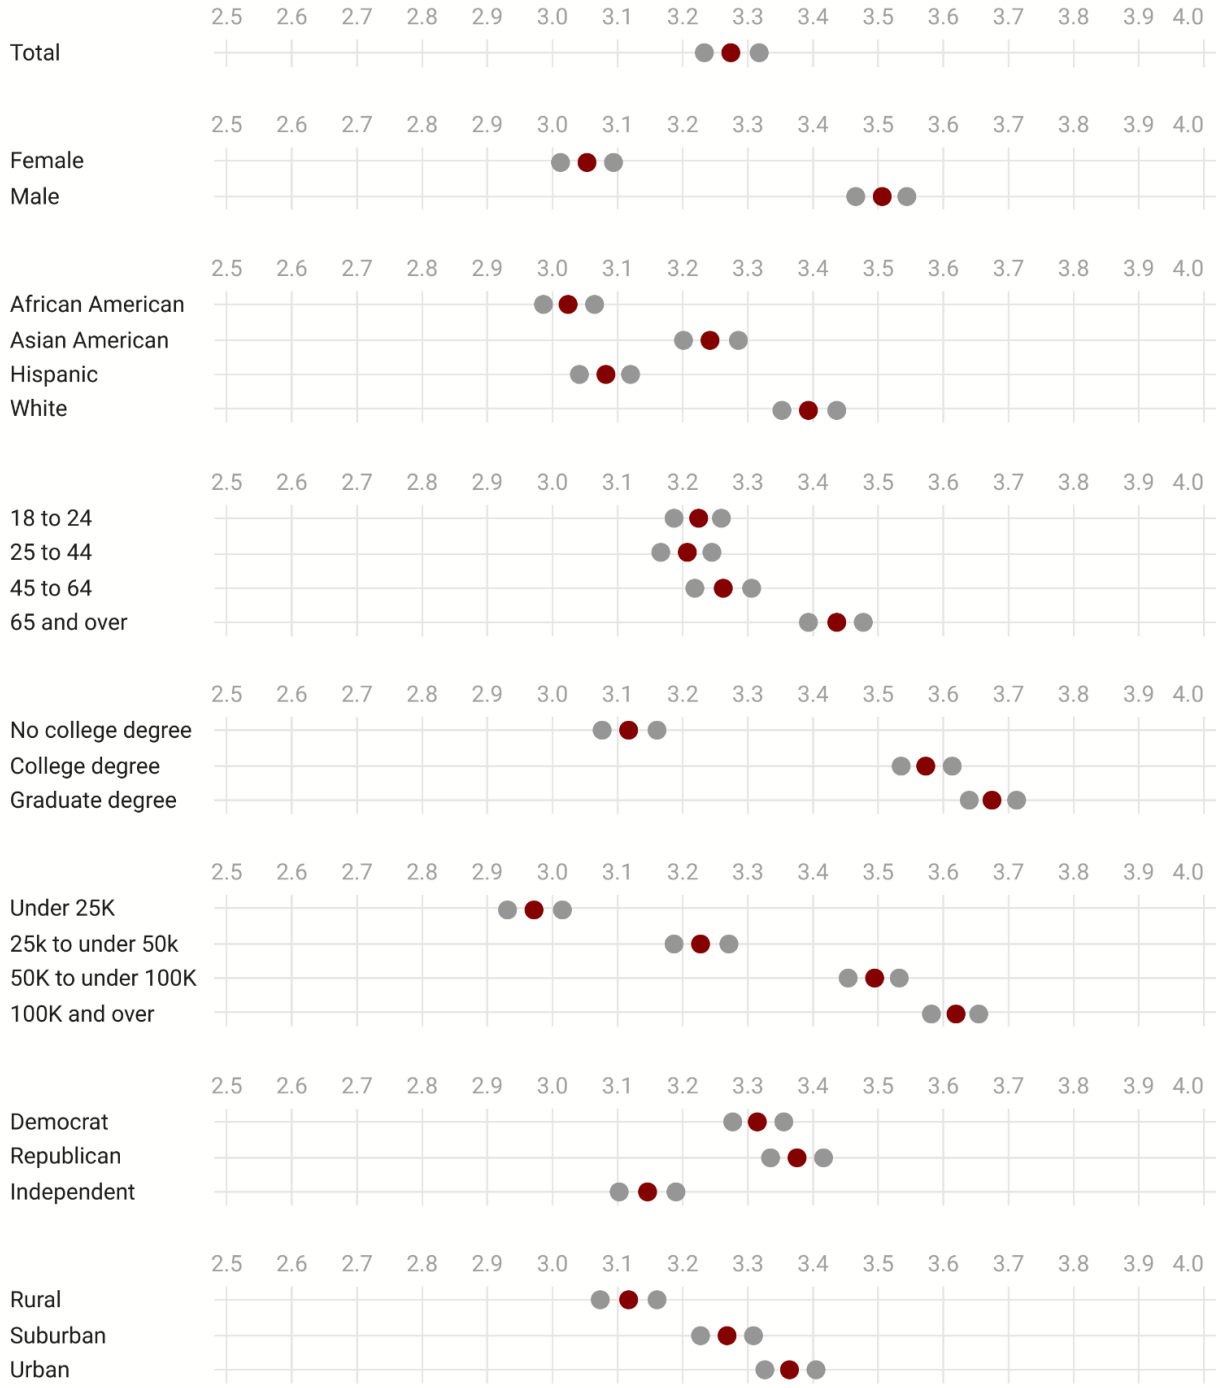

National AI Opinion Monitor (NAIOM). National sample, December 2024, N = 4767. More information at www.naiom.net.

Figure 10.

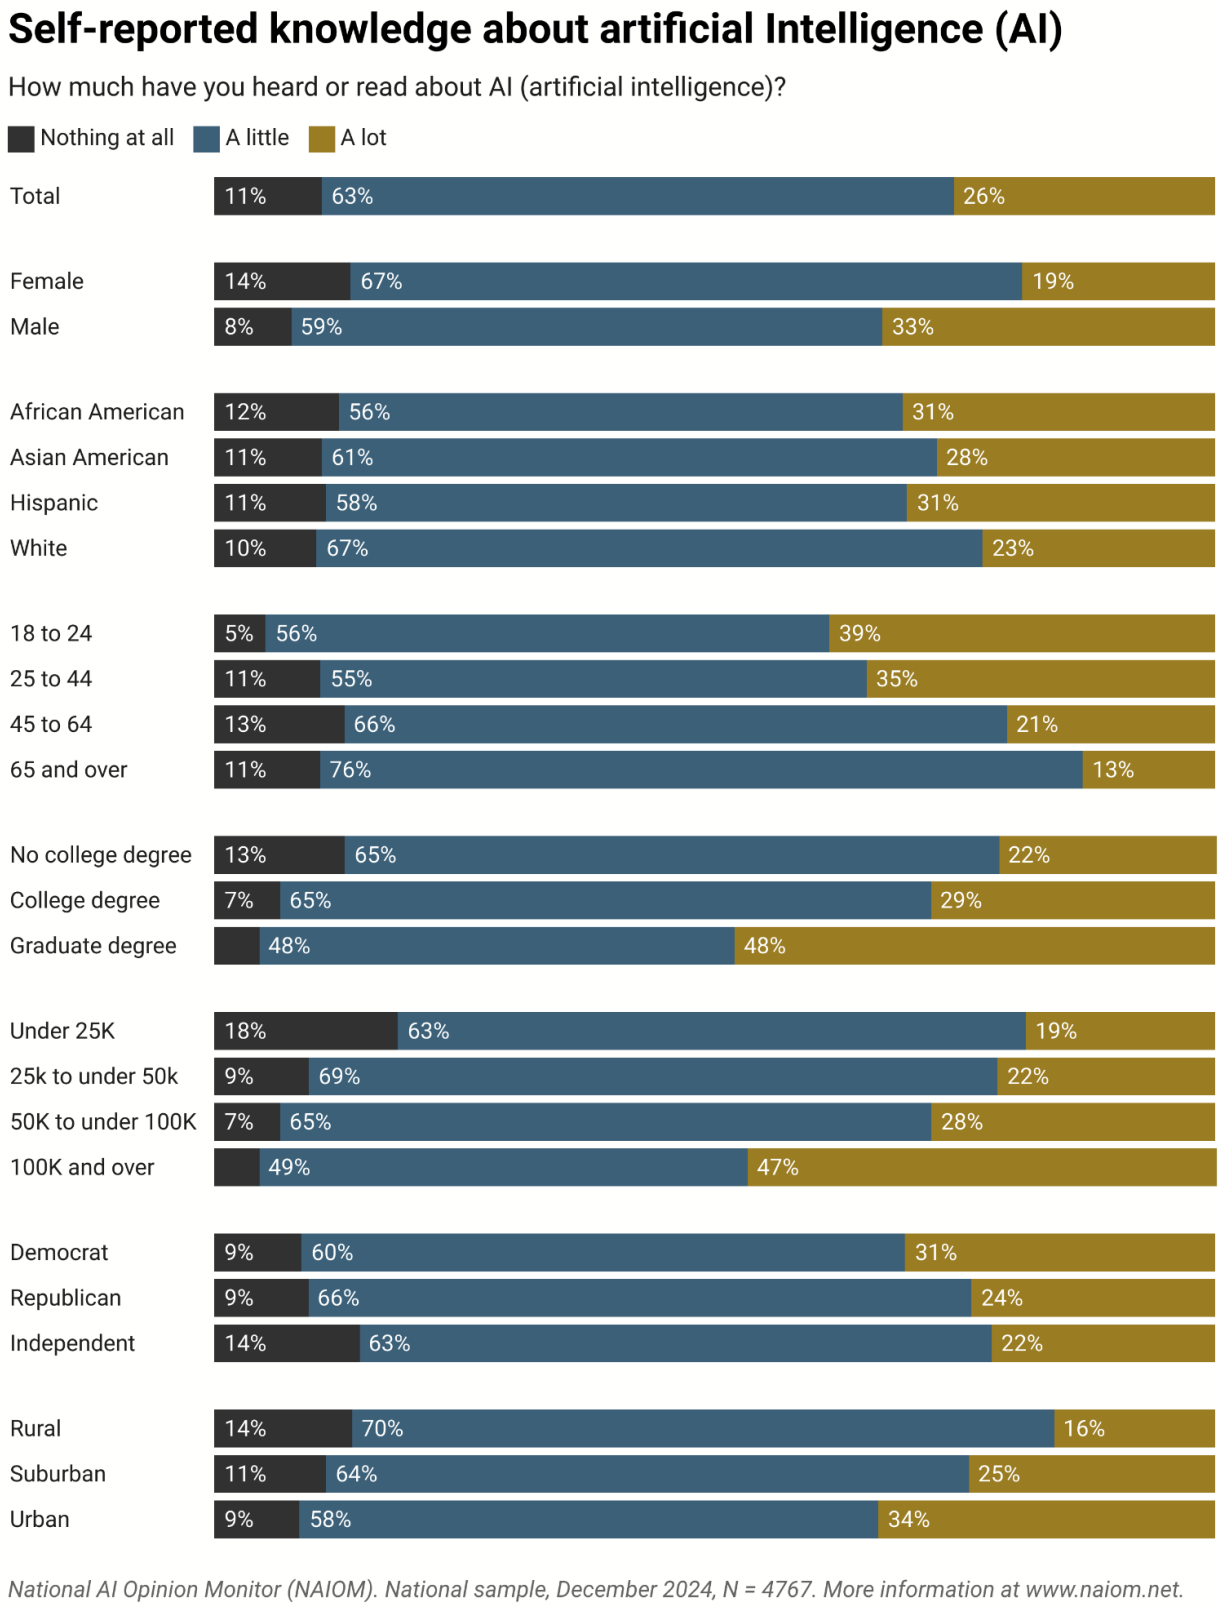
Self-reported knowledge about artificial Intelligence (AI)
How much have you heard or read about AI (artificial intelligence)?
Nothing at all
A little
A lot
Total 11% 63% 26%
Female 14% 67% 19%
Male 8% 59% 33%
African American 12% 56% 31%
Asian American 11% 61% 28%
Hispanic 11% 58% 31%
White 10% 67% 23%
18 to 24 5% 56% 39%
25 to 44 11% 55% 35%
45 to 64 13% 66% 21%
65 and over 11% 76% 13%
No college degree 13% 65% 22%
College degree 7% 65% 29%
Graduate degree 48% 48%
Under 25K 18% 63% 19%
25k to under 50k 9% 69% 22%
50K to under 100K 7% 65% 28%
100K and over 49% 47%
Democrat 9% 60% 31%
Republican 9% 66% 24%
Independent 14% 63% 22%
Rural 14% 70% 16%
Suburban 11% 64% 25%
Urban 9% 58% 34%
National AI Opinion Monitor (NAIOM). National sample, December 2024, N = 4767. More information at www.naiom.net.

Figure 11.

We also examined the average objective knowledge score for different demographics (Figure 10). One unusual pattern we see is that the average score for respondents over 65 is slightly higher than that for other age groups. The slight positive link with age, however, disappears in models where we use all demographic variables together to predict knowledge. The most important factors in those models are income, education, and gender, while age group does not significantly contribute to predicting knowledge.

To assess self-reported familiarity with artificial intelligence, respondents were asked how much they had heard or read about AI (Figure 11). Among all respondents, 26% reported having heard “a lot” about AI, while 63% had heard “a little,” and 11% had heard “nothing at all.”

Men were more likely than women to report higher familiarity, with 33% of men saying they had heard “a lot,” compared to 19% of women. Across racial and ethnic groups, African Americans, Asian Americans, and Hispanics reported similar levels of high familiarity (31%, 28%, and 31%, respectively), while only 23% of Whites reported having heard “a lot” about AI. Younger respondents aged 18-24 were the most likely to report high familiarity, with 39% saying they had heard “a lot,” while only 13% of those aged 65 and older reported the same. Education was also strongly linked to self-reported familiarity: 48% of graduate degree holders said they had heard “a lot,” compared to 22% of those without a college degree. Higher income was similarly associated with familiarity, as 47% of those earning over $100K reported hearing “a lot,” compared to 19% of those earning under $25K. Finally, urban residents reported higher familiarity (34% saying “a lot”) than suburban (25%) or rural (16%) residents.

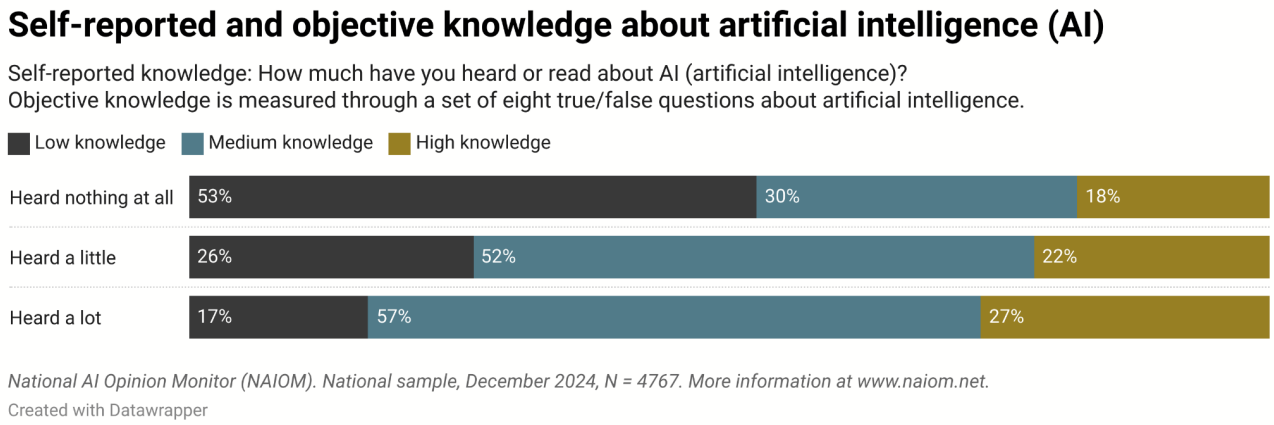

Figure 12.

We further found a linear relationship between self-reported familiarity and actual understanding of AI. Among respondents who reported hearing "nothing at all" about AI, 53% were classified as having low knowledge, while only 18% were in the high knowledge group (Figure 12). Conversely, among those who reported hearing "a lot" about AI, only

17% were classified as having low knowledge, while 27% were in the high knowledge category. These results suggest that self-reported familiarity is associated with actual knowledge but does not perfectly predict it.

Appendix A: Survey Questions

How much confidence, if any, do you have in each of the following to act in the best interest of the public?

- A great deal
- A fair amount
- Not too much
- None at all

How much do you trust companies to use artificial intelligence responsibly?

- Not at all
- Not much
- Some
- A lot

Generally speaking, would you say that most people can be trusted, or that you can't be too careful in dealing with people? On a scale from 1 to 10 (1 = "you can't be too careful" and 10 = "most people can be trusted"), would you say that…

1 2 3 4 5 6 7 8 9 10

Generally speaking, would you say that most machines can be trusted, or that you can't be too careful in dealing with machines? On a scale from 1 to 10 (1 = "you can't be too careful" and 10 = "most machines can be trusted"), would you say that…

1 2 3 4 5 6 7 8 9 10

How much would you trust news or information produced by journalists from mainstream media?

- A lot
- Some
- Not too much
- Not at all

How much would you trust news or information that was written by AI (artificial intelligence)?

- A lot
- Some
- Not too much
- Not at all

How confident are you that you can recognize whether a news story has been written by a human or AI (artificial intelligence)?

- Very confident
- Somewhat confident
- Not very confident
- Not confident at all

How much have you heard or read about AI (artificial intelligence)?

- A lot
- A little
- Nothing at all

Below are some statements about AI. To the best of your knowledge, how accurate is each of those statements? [Randomize statement order]

- AI systems can always explain how they reached a decision
- AI systems work based on specific instructions for every scenario
- AI systems learn by analyzing data and finding patterns
- AI systems make decisions by following predefined rules
- AI have thoughts and emotions like humans do

- AI systems are objective and unbiased
- AI can be used to make art and music
- AI can be used to help diagnose diseases
  - Accurate
  - Inaccurate
  - Not sure